

## New York City Airbnb Exploratory Data Analysis



### A) Background

#### Context

*Since 2008, guests and hosts have used Airbnb to expand on traveling possibilities and present more unique, personalized way of experiencing the world. This dataset describes the listing activity and metrics in NYC, NY for 2019.*

#### Content

*This data file includes all needed information to find out more about hosts, geographical availability, necessary metrics to make predictions and draw conclusions.*

#### Acknowledgements

*This public dataset is part of Airbnb, I pulled this data from Kaggle*  
<https://www.kaggle.com/dgomonov/new-york-city-airbnb-open-data>

#### Some Questions to answer :

- 1) What can we learn about different hosts and areas?
- 2) What can we learn from predictions? (ex: locations, prices, reviews, etc)
- 3) Which hosts are the busiest and why?
- 4) Is there any noticeable difference of traffic among different areas and what could be the reason for it?

## B) Import the Libraries :

```
import pandas as pd
import numpy as np
import seaborn as sns

import matplotlib.pyplot as plt
from matplotlib import rcParams

import warnings
warnings.filterwarnings('ignore')
```

## C) Get the Data (Read the .csv file) :

```
data = pd.read_csv('/content/AB_NYC_2019.csv')
```

## D) Let's explore the data :

- i. Data Overview
- ii. No. of Records & Features
- iii. Summary & info
- iv. Conclusion

### i. Data Overview :

*# Top 5 records :*

```
data.head(5)
```

	id	...	availability_365
0	2539	...	365
1	2595	...	355
2	3647	...	365
3	3831	...	194
4	5022	...	0

[5 rows x 16 columns]

### ii. No. of Records & Features :

*# Let's know our features:*

```
print('SHAPE OF OUR DATASET :',data.shape)
```

```
print('\n WHAT ARE OUR FEATURES ? : \n',data.columns)
```

SHAPE OF OUR DATASET : (48895, 16)

WHAT ARE OUR FEATURES ? :

```
Index(['id', 'name', 'host_id', 'host_name', 'neighbourhood_group',
      'neighbourhood', 'latitude', 'longitude', 'room_type', 'price',
      'minimum_nights', 'number_of_reviews', 'last_review',
      'reviews_per_month', 'calculated_host_listings_count',
      'availability_365'],
      dtype='object')
```

### OBSERVATION :

1) There are around 48895 records in our dataset.

2) We have 16 features :

id, name, host\_id, host\_name, neighbourhood\_group, neighbourhood,  
latitude, longitude, room\_type, price, minimum\_nights,  
number\_of\_reviews, last\_review, reviews\_per\_month,  
calculated\_host\_listings\_count, availability\_365

### iii. Descriptive Summary

```
data.describe()
```

	id	...	availability_365
count	4.889500e+04	...	48895.000000
mean	1.901714e+07	...	112.781327
std	1.098311e+07	...	131.622289
min	2.539000e+03	...	0.000000
25%	9.471945e+06	...	0.000000
50%	1.967728e+07	...	45.000000
75%	2.915218e+07	...	227.000000
max	3.648724e+07	...	365.000000

[8 rows x 10 columns]

# INFO

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 48895 entries, 0 to 48894
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	id	48895 non-null	int64
1	name	48879 non-null	object
2	host_id	48895 non-null	int64
3	host_name	48874 non-null	object
4	neighbourhood_group	48895 non-null	object
5	neighbourhood	48895 non-null	object

```

6 latitude 48895 non-null float64
7 longitude 48895 non-null float64
8 room_type 48895 non-null object
9 price 48895 non-null int64
10 minimum_nights 48895 non-null int64
11 number_of_reviews 48895 non-null int64
12 last_review 38843 non-null object
13 reviews_per_month 38843 non-null float64
14 calculated_host_listings_count 48895 non-null int64
15 availability_365 48895 non-null int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB

```

#### iv. Conclusion :

- We have total 16 features, out of which "*last\_review*", "*name*", "*host\_name*" are least important in my opinion. (So we'll ignore it)
- Factors/Metrics to look into
  - a) price    b) number\_of\_reviews    c) reviews\_per\_month    d) availability\_365    e) calculated\_host\_listings\_count
- We can group the data based on
  - a) host\_id, host\_name    b) neighbourhood\_group    c) neighbourhood    d) room\_type
- We can utilise Longitude & Latitude values for geospatial visualization
- We have some features with *null* values.

## E) Data Wrangling :

- Duplicate Values
- Missing Values

#### i. Duplicate Values :

# For [id] :

```

print("Total no. of id values :", data['id'].count())
print("No. of unique id values :", len(data['id'].unique()))

```

# For [name] :

```

print("\nTotal no. of name values : ", data['name'].count())
print("No. of unique name values : ", len(data['name'].unique()))

```

```

# For [host_id] :

print("\nTotal no. of host_id values :", data['host_id'].count())
print("No. of unique host_id values :", len(data['host_id'].unique()))

# For [host_name] :

print("\nTotal no. of host_name values : ", data['host_name'].count())
print("No. of unique host_name values : ", len(data['host_name'].unique()))

Total no. of id values : 48895
No. of unique id values : 48895

Total no. of name values : 48879
No. of unique name values : 47906

Total no. of host_id values : 48895
No. of unique host_id values : 37457

Total no. of host_name values : 48874
No. of unique host_name values : 11453

```

### Observation :

- As we can see, we have some duplicate records when it comes to 'name' & 'host\_name'
- We don't need these 3 features:
  - a) name      b) host\_name      c) last\_review

So, we'll drop these features.
- Also, we even have null values in this features & they don't seem really helpful. So, let's drop them.

### # Drop these 3 features in our DataFrame

*# Drop the features from the dataframe*

```
cleaned_data = data.drop(['name', 'host_name', 'last_review'], axis=1)
```

```
cleaned_data
```

```

      id  host_id  ...  calculated_host_listings_count
availability_365
0      2539    2787  ...                             6
365

```

```

1          2595      2845 ...          2
355
2          3647      4632 ...          1
365
3          3831      4869 ...          1
194
4          5022      7192 ...          1
0
...          ...      ... ...          ...
...
48890  36484665    8232441 ...          2
9
48891  36485057    6570630 ...          2
36
48892  36485431    23492952 ...          1
27
48893  36485609    30985759 ...          6
2
48894  36487245    68119814 ...          1
23

```

[48895 rows x 13 columns]

## ii. Missing Values :

*# Find the missing / NaN values :*

```
print(cleaned_data.isnull().sum())
```

```

id          0
host_id     0
neighbourhood_group  0
neighbourhood  0
latitude     0
longitude     0
room_type    0
price        0
minimum_nights  0
number_of_reviews  0
reviews_per_month 10052
calculated_host_listings_count  0
availability_365  0
dtype: int64

```

## Observation :

- There's just 1 feature with huge no. of missing values i.e. *reviews\_per\_month*
- Let's replace this missing values with '0'

```
# Fill the missing values :
```

```
cleaned_data.fillna(value = 0,inplace = True)
```

```
# New Cleaned Data
```

```
cleaned_data
```

```
      id  host_id  ... calculated_host_listings_count
availability_365
0      2539     2787  ...                             6
365
1      2595     2845  ...                             2
355
2      3647     4632  ...                             1
365
3      3831     4869  ...                             1
194
4      5022     7192  ...                             1
0
...      ...      ...  ...
...
48890  36484665   8232441  ...                             2
9
48891  36485057   6570630  ...                             2
36
48892  36485431  23492952  ...                             1
27
48893  36485609  30985759  ...                             6
2
48894  36487245  68119814  ...                             1
23
```

```
[48895 rows x 13 columns]
```

## F) Exploratory Data Analysis :

- i. Host
- ii. price
- iii. room\_type
- iv. Area (neighborhood, neighborhood\_group)
- v. availability\_365

### i. Host

- Who are the busiest hosts?
- Which areas do these hosts cater to?

```
cleaned_data['host_id'].value_counts()
```

```

219517861    327
107434423    232
30283594     121
137358866    103
12243051      96
...
1641589      1
4070519      1
208106618    1
235939247    1
1288080      1
Name: host_id, Length: 37457, dtype: int64

```

### ***Q1 . Who is the most busiest host in New York City?***

```
cleaned_data['host_id'].mode()
```

```

0    219517861
dtype: int64

```

#### **Observation :**

- '219517861' is the most busiest host
- '107434423', '30283594', '137358866', '12243051' are the other top hosts

```
top_hosts = [219517861, 107434423, 30283594, 137358866, 12243051]
```

```
top_hosts
```

```
[219517861, 107434423, 30283594, 137358866, 12243051]
```

```
print('Top Hosts Name : \n')
```

```

for i in top_hosts:
    print(data[data['host_id'] == i].host_name.unique())

```

```
Top Hosts Name :
```

```

['Sonder (NYC)']
['Blueground']
['Kara']
['Kazuya']
['Sonder']

```

### ***Q2. What are the areas these busiest hosts belong to ?***

```
features = ['host_id', 'neighbourhood_group']
```

```
busiest_hosts = cleaned_data[cleaned_data['host_id'].isin(top_hosts)]
```

```
busiest_hosts_areas = busiest_hosts.loc[:,
```



```
['host_id', 'neighbourhood_group']]
busiest_hosts_areas.drop_duplicates(inplace = True)
```

busiest\_hosts\_areas

	host_id	neighbourhood_group
9740	30283594	Manhattan
26137	107434423	Manhattan
30637	12243051	Manhattan
32718	137358866	Manhattan
33268	137358866	Queens
36698	137358866	Brooklyn
38293	219517861	Manhattan
39275	107434423	Brooklyn

Observation :

- *'Sonder (NYC) ', 'Blueground ', 'Kara ', 'Kazuya ', 'Sonder '* are the busiest hosts.
- These busiest hosts mostly belong to *Manhattan, Brooklyn & Queens*.

## ii. Price

- Price distribution
- Median Price in \$
- Price Statistic
- At what price do the busiest/top hosts offer their services?
- Prices across different regions

```
prices_df = cleaned_data.loc[:,
['id', 'host_id', 'neighbourhood_group', 'price']]
```


prices\_df

	id	host_id	neighbourhood_group	price
0	2539	2787	Brooklyn	149
1	2595	2845	Manhattan	225
2	3647	4632	Manhattan	150
3	3831	4869	Brooklyn	89
4	5022	7192	Manhattan	80
...	...	...	...	...
48890	36484665	8232441	Brooklyn	70
48891	36485057	6570630	Brooklyn	40
48892	36485431	23492952	Manhattan	115
48893	36485609	30985759	Manhattan	55
48894	36487245	68119814	Manhattan	90

[48895 rows x 4 columns]

# Price Distribution :

```
sns.distplot(prices_df['price'])  
plt.xlim(0, 1000)  
(0.0, 1000.0)
```



**# Median Price in \$ :**


```
print('Median Price ',prices_df['price'].median())
```

Median Price 106.0

**# Price Statistic :**

```
sns.boxplot(x = prices_df['price'])  
plt.xlim(0,1000)
```

```
(0.0, 1000.0)
```



*# Percentile values :*

```
print('\n25th percentile : ',prices_df['price'].quantile(q=0.25))
print('\n50th percentile : ',prices_df['price'].quantile(q=0.5))
print('\n75th percentile : ',prices_df['price'].quantile(q=0.75))
print('\n99th percentile : ',prices_df['price'].quantile(q=0.99))
```

25th percentile : 69.0

50th percentile : 106.0

75th percentile : 175.0

99th percentile : 799.0

**Observation :**

- Mostly the prices lies between \$ 69 & \$ 175
- Median price across all of New York City is \$ 106
- Also, our price data contains huge amount of outliers (which is why we used Median in place of Averages/Mean value)

**# At what price do the busiest/top hosts offer their services?**

```
busiest_hosts_prices = busiest_hosts.loc[:,['host_id','price']]
```

```
busiest_hosts_prices
```


	host_id	price
9740	30283594	169
10075	30283594	135
10335	30283594	369
10398	30283594	335
10490	30283594	129
...	...	...
48723	107434423	316
48724	107434423	385
48725	107434423	267
48726	107434423	278
48727	107434423	365

```
[879 rows x 2 columns]
```

```
sns.distplot(busiest_hosts_prices['price'])
```

```
plt.xlim(0, 1000)
```


```
(0.0, 1000.0)
```



```
sns.boxplot(x = busiest_hosts_prices['price'])
```

```
plt.xlim(0,450)
```

```
(0.0, 450.0)
```



*# Percentile values :*

```
print('\n25th
percentile :',busiest_hosts_prices['price'].quantile(q=0.25))
print('\n50th
percentile :',busiest_hosts_prices['price'].quantile(q=0.5))
print('\n75th
percentile :',busiest_hosts_prices['price'].quantile(q=0.75))
print('\n99th
percentile :',busiest_hosts_prices['price'].quantile(q=0.99))
```

25th percentile : 185.0

50th percentile : 239.0

75th percentile : 289.0

99th percentile : 699.0


### **Observation :**

- Turns out that the Busiest host's in New York City cost a little more expensive than others. i.e. between \$ 185 & \$ 289

**# Prices across different regions :**

```
sns.boxplot(x="neighbourhood_group", y="price", data = prices_df)
plt.ylim(0,500)
```

(0.0, 500.0)



#### Observation :

- Manhattan is the most expensive area. (highest price)
- Manhattan & Brooklyn turns out to be the most expensive areas.
- Rest, { Queens, Staten Island, Bronx } have similar prices and are cheaper than Manhattan & Brooklyn.

#### iii. Room Type :

- Which room types are more popular?
- Which room types are more cheaper than others?

```
roomtype_df = cleaned_data.loc[:,['id', 'room_type', 'price']]
```

```
roomtype_df
```

	id	room_type	price
0	2539	Private room	149
1	2595	Entire home/apt	225
2	3647	Private room	150
3	3831	Entire home/apt	89
4	5022	Entire home/apt	80
...	...	...	...
48890	36484665	Private room	70
48891	36485057	Private room	40
48892	36485431	Entire home/apt	115
48893	36485609	Shared room	55

```
48894 36487245 Private room 90
```

```
[48895 rows x 3 columns]
```

```
# Popularity of different room_types :
```


```
print('Distribution of Room Types : \n\n',roomtype_df['room_type'].value_counts())
```

```
Distribution of Room Types :
```

```
Entire home/apt    25409
Private room       22326
Shared room        1160
Name: room_type, dtype: int64
```

```
# Visualize the above result in form of pie-chart :
```

```
count = [25409,22326,1160]
labels = ['Entire home/apt','Private room','Shared room']
colors = sns.color_palette('pastel')
plt.pie(count, labels=labels,colors = colors, autopct = '%0.0f%%')
plt.show()
```

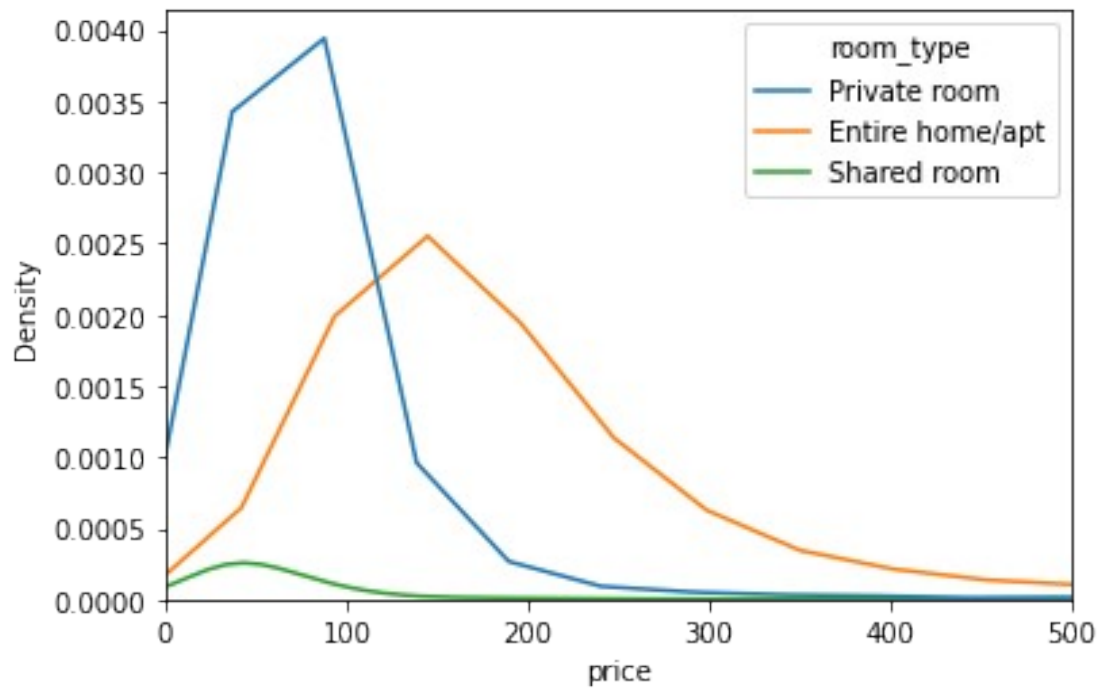


**Observation :**

- As we can clearly see, 'Entire home/apt' and 'Private room' are the more popular room types.
- Shared room is the least preferred room type contributing to just 2 % of all bookings.


```
# Price of different room_types
```

```
sns.kdeplot(data = roomtype_df, x="price", hue="room_type")  
plt.xlim(0,500)  
(0.0, 500.0)
```



```
sns.boxplot(x="room_type", y="price", data = roomtype_df)  
plt.ylim(0,500)  
(0.0, 500.0)
```





#### Observation :

- "Entire Home/Apt" is the most expensive room type.
- "Private Room" & "Shared Room" are almost comparable in terms of Price.
- "Shared Room" is the cheapest room type.

#### iv. Area (neighbourhood\_group, neighbourhood):

- Which are the most busiest/popular areas in New York City?

```
neighbour_group = cleaned_data.groupby(by='neighbourhood_group')
```

```
neighbour_group
```

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x7fcfce117f50>
```

```
neighbour_group.describe()
```

		id		...
availability_365		count	mean	std
50%	75%	max		
neighbourhood_group				...
Bronx		1091.0	2.273492e+07	1.023402e+07
148.0	313.5	365.0		...
Brooklyn		20104.0	1.825685e+07	1.083320e+07
28.0	188.0	365.0		...
Manhattan		21661.0	1.877494e+07	1.116793e+07

```

36.0  230.0  365.0
Queens                    5666.0  2.175500e+07  1.037687e+07  ...
98.0  286.0  365.0
Staten Island             373.0  2.159747e+07  1.039310e+07  ...
219.0  333.0  365.0

```

```
[5 rows x 80 columns]
```

```
neighbour_group['id'].count()
```

```

neighbourhood_group
Bronx                1091
Brooklyn             20104
Manhattan            21661
Queens               5666
Staten Island        373
Name: id, dtype: int64

```

```
g = dict(neighbour_group['id'].count())
```

```

keys = list(g.keys())
values = list(g.values())


```

```

plt.title('Bookings across different Neighbourhood_groups')
colors = sns.color_palette('pastel')
plt.pie(values, labels=keys, colors = colors, autopct = '%0.0f%%')
plt.show()

```

Bookings across different Neighbourhood\_groups



**Observation :**


- Manhattan & Brooklyn are the busiest areas (neighbourhood\_group). These two contributes to around 85 % of the total traffic

#### vi. Availability\_365 :

# *Availability across all hosts :*


```
sns.distplot(cleaned_data['availability_365'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcfca71ded0>
```



```
sns.boxplot(x=cleaned_data['availability_365'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcfcd275690>
```



*# Percentile values :*

```
print('\n25th
percentile : ',cleaned_data['availability_365'].quantile(q=0.25))
print('\n50th
percentile : ',cleaned_data['availability_365'].quantile(q=0.5))
print('\n75th
percentile : ',cleaned_data['availability_365'].quantile(q=0.75))
print('\n99th
percentile : ',cleaned_data['availability_365'].quantile(q=0.99))
```

25th percentile : 0.0

50th percentile : 45.0

75th percentile : 227.0


99th percentile : 365.0

### **Observation :**


- Most of the Airbnb host's offer their services for over 0 to 227 days in a year.

**# Availability across the busiest hosts :**

```
sns.distplot(busiest_hosts['availability_365'])
rcParams['figure.figsize'] = 7,5
```



```
sns.boxplot(x=busiest_hosts['availability_365'])
<matplotlib.axes._subplots.AxesSubplot at 0x7fcfcb003d10>
```



*# Percentile values :*

```
print('\n25th
```

```

percentile :',busiest_hosts['availability_365'].quantile(q=0.25))
print('\n50th
percentile :',busiest_hosts['availability_365'].quantile(q=0.5))
print('\n75th
percentile :',busiest_hosts['availability_365'].quantile(q=0.75))
print('\n99th
percentile :',busiest_hosts['availability_365'].quantile(q=0.99))

```

25th percentile : 235.0

50th percentile : 311.0

75th percentile : 339.0

99th percentile : 365.0

### Observation :

- The busiest Airbnb host's offer their services for over 235 to 339 days in a year.

## G) Correlation between Features :

```
new_df = cleaned_data.drop(['longitude','latitude'], axis=1)
```

```
corr_df = new_df.corr()
```

```

sns.heatmap(corr_df, cmap="YlGnBu", annot=True)
rcParams['figure.figsize'] = 18,8
plt.show()

```



## H) Conclusion :

- Our Dataset contained 48,895 Airbnb booking records around New York City for the year of 2019.
- We were provided with 16 features, out of which we used some of the features for the purpose of analysis.
- '*Sonder (NYC)* ', '*Blueground* ', '*Kara* ', '*Kazuya* ', '*Sonder* ' are the busiest hosts. These hosts cater around *Manhattan & Brooklyn* areas.
- Manhattan & Brooklyn are the most popular areas in New York City with the busiest traffic. These areas are also the most expensive ones. (contributing to 85% of the total traffic)
- Manhattan is the most expensive area.
- 'Entire home/Apt' and 'Private room' are the most preferred room categories.
- 'Entire home/Apt' is the most expensive room category.
- 'Shared rooms' are hardly booked(only 2% out of total bookings), even though it is the cheapest room category.
- While most of the Airbnb host's offer their services for over 0 to 227 days in a year. The busiest Airbnb hosts offer their services from 235 to 339 days in a year. (which is so much more than other hosts)

## I) Recommendation :

- Customers surely seems to value '**quality experience**' over '**money**'. So, Airbnb should advise their hosts to prioitize on providing a high quality experience & then later align their prices.
- **Manhattan, Brooklyn** and **Queens** receives the best traffic among all the neighbourhood areas. So, Airbnb should find more hosts catering to these areas.
- '**Shared Rooms**' don't get much customer traffic. So it would be a wise idea to remove this room category from the Airbnb services.